

Statistical methods for bioinformatics

GAM and Trees

Cedric Lood

May 15, 2016

1 Applied exercises

1.1 Question 10

Libraries and definition of the training and test sets used for the analysis:

```
library(ISLR)
library(boot)
library(ggplot2)
library(leaps)
library(gam)

attach(College)
set.seed(1)
train <- sample(c(TRUE,FALSE), nrow(College), rep=TRUE)
test <- (!train)
```

1.1.1 Part a

The analysis reveals a forward selection with 6 variables retained in the model:

```
model.fwd <- regsubsets(Outstate~.,data=College[train,], nvmax=17,method="forward")
test.mat <- model.matrix(Outstate~., data=College[test,])
```

```
val.errors <- rep(NA, 17)
for(i in 1:17){
  coefi <- coef(model.fwd, id=i)
  pred <- test.mat[,names(coefi)] %*% coefi
  val.errors[i] <- mean((College$Outstate[test]-pred)^2)
}
val.errors
which.min(val.errors)
coef(model.fwd, 6)
```

```
> val.errors
[1] 10734659 7647452 6468424 5434378 4948421 4650921 4734848 4718857
[9] 4690641 4735896 4694172 4877469 4754741 4690085 4711840 4700313
```

```
[17] 4700869
> which.min(val.errors)
[1] 6
> coef(model.fwd, 6)
      (Intercept)      PrivateYes      Room.Board      Terminal      perc.alumni
-4227.6797221    2778.7052614      0.8106532    49.5653264    38.3960635
      Expend      Grad.Rate
    0.2616141    26.3975500
```

1.1.2 Part b

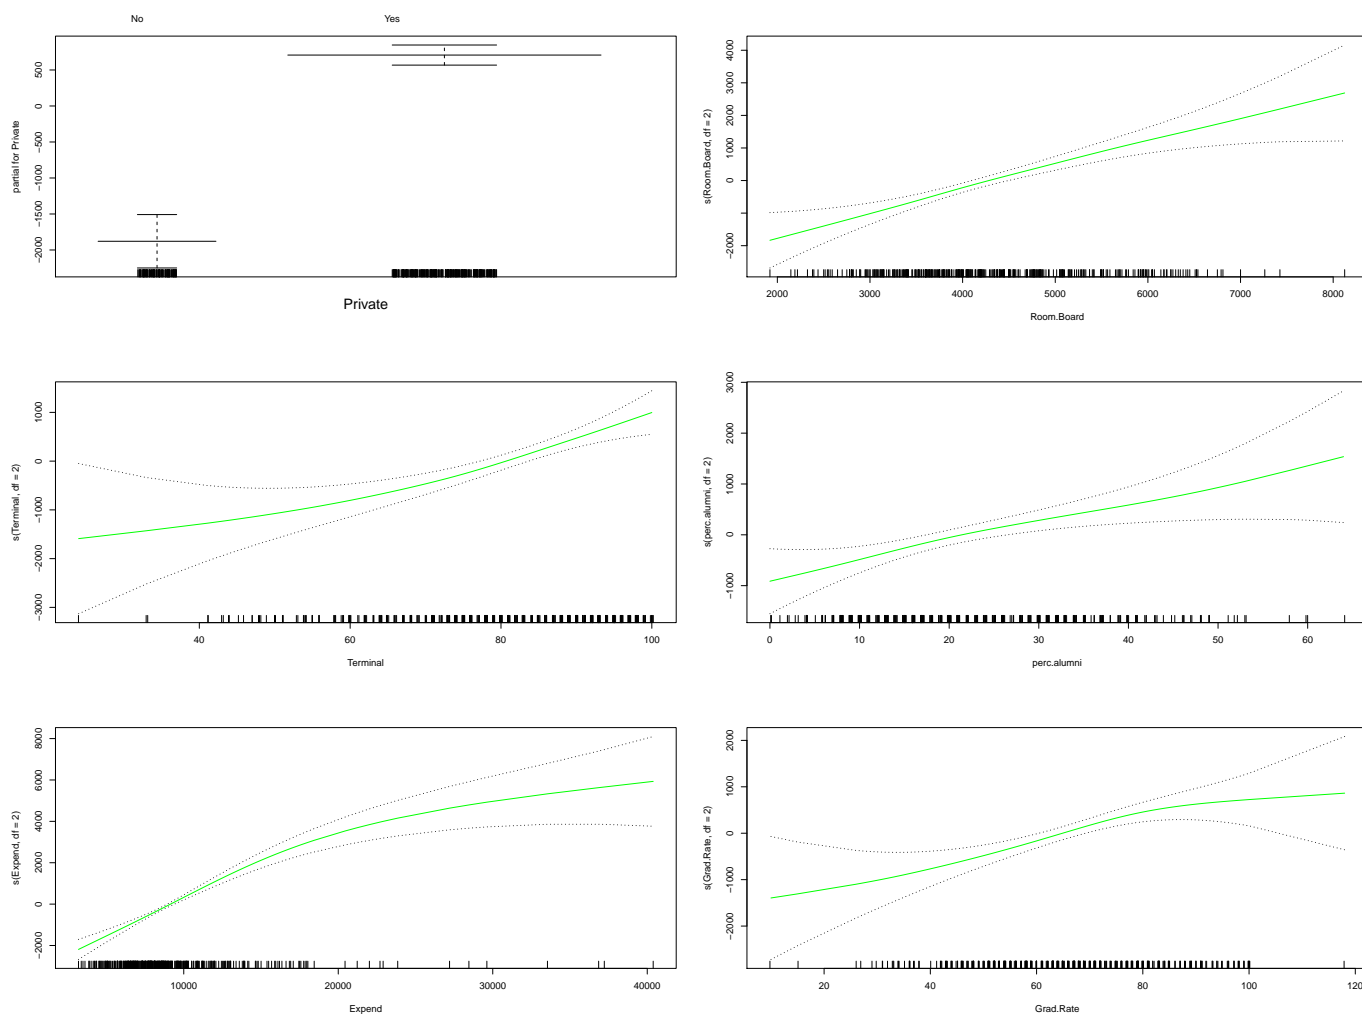
For this section, there are potentially many different GAM that could be produced. For example, one could produce different functionals consisting of smoothing splines for each of the predictors, with different combinations of degrees of freedom. Here are a few, the first one being akin to a multiple linear regression:

```
model.gam <- gam(Outstate~Private+Room.Board+Terminal+perc.alumni+Expend+Grad.Rate,
                 data=College[train,])
pdf("gam_trees_s.pdf", width=16, height=12)
par(mfrow = c(3, 2))
plot.gam(model.gam,se=TRUE,col="green")
dev.off()
```

```
model.gam.s2 <- gam(Outstate~Private+s(Room.Board,df=2)+s(Terminal,df=2)+
                   s(perc.alumni,df=2)+s(Expend,df=2)+s(Grad.Rate,df=2),
                   data=College[train,])
pdf("gam_trees_s_2.pdf", width=16, height=12)
par(mfrow = c(3, 2))
plot.gam(model.gam.s,se=TRUE,col="green")
dev.off()
```

```
model.gam.s.3 <- gam(Outstate~Private+s(Room.Board,df=3)+s(Terminal,df=3)+
                   s(perc.alumni,df=3)+s(Expend,df=3)+s(Grad.Rate,df=3),
                   data=College[train,])
pdf("gam_trees_s_3.pdf", width=16, height=12)
par(mfrow = c(3, 2))
plot.gam(model.gam.s.3,se=TRUE,col="green")
dev.off()
```

The graphics for the other 2 models can be found in the annex. Here is the one for the smoothing splines with 2 degrees of freedom:



1.1.3 Part c

This is the evaluation of the 3 different models for the test sets, followed by the results.

```
gam.pred <- predict(model.gam,College[test,])
gam.err <- mean((College[test,]$Outstate - gam.pred)^2)

gam.pred.s2 <- predict(model.gam.s2,College[test,])
gam.err.s2 <- mean((College[test,]$Outstate - gam.pred.s2)^2)

gam.pred.s3 <- predict(model.gam.s.3,College[test,])
gam.err.s3 <- mean((College[test,]$Outstate - gam.pred.s3)^2)

> gam.err
[1] 4650921
> gam.err.s2
[1] 3952853
> gam.err.s3
[1] 3784922
```

1.1.4 Part d

Evidence of non-linear relationship between variables can be investigated using the summary function on the GAM. The function provides the user with a non-parametric ANOVA table. Below is the output of the command, which seem to indicate a non-linear relationship between the predictor “Expend” and the dependent value.

```
> summary(model.gam.s2)
```

```
Call: gam(formula = Outstate ~ Private + s(Room.Board, df = 2) + s(Terminal,
      df = 2) + s(perc.alumni, df = 2) + s(Expend, df = 2) + s(Grad.Rate,
      df = 2), data = College[train, ])
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-7045.4	-1145.4	84.8	1184.1	5047.0

(Dispersion Parameter for gaussian family taken to be 3452101)

Null Deviance: 6006262152 on 405 degrees of freedom

Residual Deviance: 1360128366 on 394.0002 degrees of freedom

AIC: 7278.121

Number of Local Scoring Iterations: 2

Anova for Parametric Effects

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
Private	1	1635387248	1635387248	473.737	< 2.2e-16 ***
s(Room.Board, df = 2)	1	1343549886	1343549886	389.198	< 2.2e-16 ***
s(Terminal, df = 2)	1	597441375	597441375	173.066	< 2.2e-16 ***
s(perc.alumni, df = 2)	1	240771844	240771844	69.746	1.161e-15 ***
s(Expend, df = 2)	1	424246993	424246993	122.895	< 2.2e-16 ***
s(Grad.Rate, df = 2)	1	63996091	63996091	18.538	2.104e-05 ***
Residuals	394	1360128366	3452101		

Signif. codes: 0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Anova for Nonparametric Effects

	Npar	Df	Npar F	Pr(F)
(Intercept)				
Private				
s(Room.Board, df = 2)	1	0.5670	0.4519	
s(Terminal, df = 2)	1	2.2148	0.1375	
s(perc.alumni, df = 2)	1	1.0512	0.3059	
s(Expend, df = 2)	1	26.1955	4.837e-07	***
s(Grad.Rate, df = 2)	1	3.5803	0.0592	.

Signif. codes: 0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

1.2 Trees Vijver

Libraries and definition of the training dataset (126 cases out of 188) used for the analysis.

```
library(glmnet)
library(tree)
library(randomForest)
library(gbm)
library(ggplot2)

load("VIJVER.Rdata")
set.seed(1)
train <- sample(1:nrow(x), 126)
data.test <- data[-train,]
meta.test <- data$meta[-train]
```

1.2.1 Performance with Ridge and Lasso

For reminder, here are the performance obtained using regularization techniques:

```
> perf.ridge
[1] 0.7580645
```

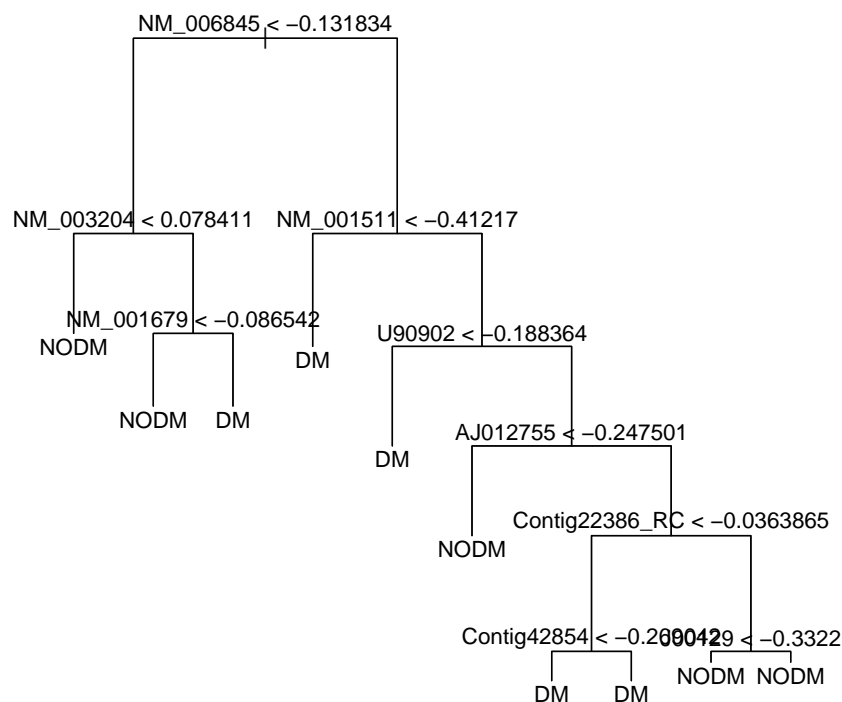
```
> perf.lasso
[1] 0.6290323
```

1.2.2 Classification tree

```
tree.data <- tree(meta~., data=data, subset=train)
tree.pred <- predict(tree.data, data.test, type="class")
table(tree.pred, meta.test)
(15+25)/62 # perf=64.5%
```

```
pdf("tree_simple.pdf")
plot(tree.data)
text(tree.data, pretty=0)
dev.off()
```

```
> table(tree.pred, meta.test)
      meta.test
tree.pred DM NODM
      DM   15   9
      NODM 13  25
> (15+25)/62
[1] 0.6451613
```



1.2.3 Classification tree pruned

Pruning the tree results in a smaller tree at a very small penalty in terms of predictive power (64% -> 63%)

```

cv.data <- cv.tree(tree.data, FUN=prune.misclass)
cv.data # best depth=4
prune.data <- prune.misclass(tree.data, best=4)
tree.pred <- predict(prune.data, data.test, type="class")
table(tree.pred, meta.test)
(13+26)/62 # perf=63%

```

```

pdf("tree_pruned.pdf")
plot(prune.data)
text(prune.data, pretty=0)
dev.off()

```

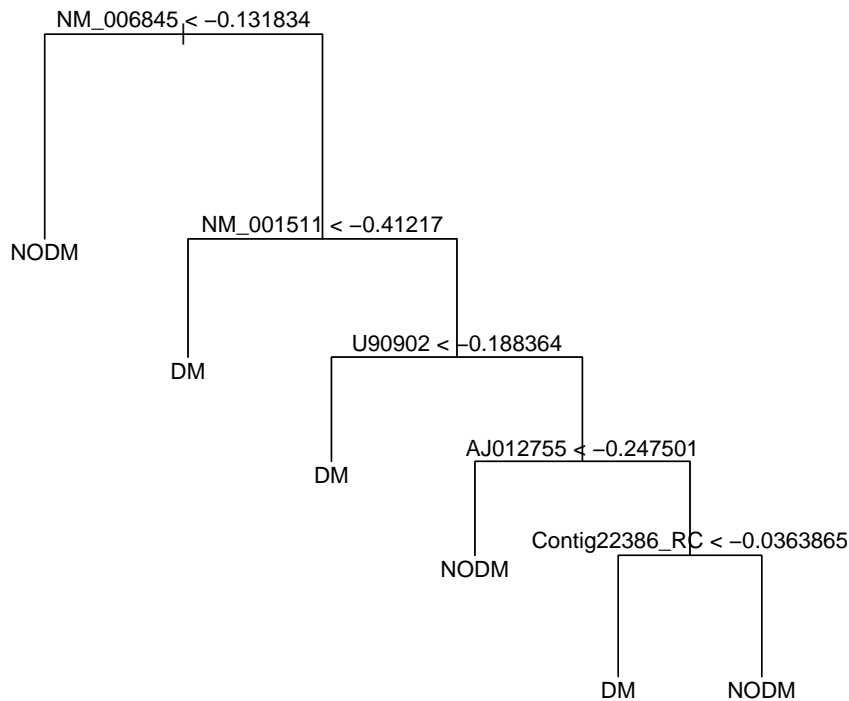
>

```

      meta.test
tree.pred DM NODM
      DM   13    8
      NODM 15   26

```

```
> (13+26)/62 # perf=63%
[1] 0.6290323
```



1.2.4 Intermezzo

For the techniques that follow, I used the variable selected by lasso to proceed with the analysis. Trying to fit any of the ensemble techniques, or boosting was not feasible on the full set of predictors (about 5000) due to computational complexity. The variable selected through lasso were:

```
> lasso.predictors
[1] "NM_000918"      "NM_003258"      "NM_004119"      "AF279865"
[5] "NM_002811"      "Contig48919_RC" "NM_003714"      "NM_006054"
[9] "NM_007267"      "AF055033"       "Contig59134_RC" "NM_001007"
```

And the definition of the datasets used for the next techniques:

```
data.lasso <- data[c("meta", predictors.lasso)]
data.test <- data.lasso[-train,]
meta.test <- data.lasso$meta[-train]
```

Of all these more advanced techniques, random forest obtained the better score in terms of classification, with 68% on the test dataset. Followed by bagging (66%) and boosting, which underperformed the simpler tree established in the previous section.

1.2.5 Bagging

```
bag.data <- randomForest(meta~., data=data.lasso, subset=train, mtry=12, importance=TRUE)
yhat.bag <- predict(bag.data, newdata=data.test)
plot(yhat.bag, meta.test)
table(yhat.bag, meta.test)
(19+22)/62
importance(bag.data)

> table(yhat.bag, meta.test)
      meta.test
yhat.bag DM  NODM
      DM   18   11
      NODM 10   23
> (18+23)/62
[1] 0.6612903
```

1.2.6 Random Forest

```
rf.data <- randomForest(meta~., data=data.lasso, subset=train, importance=TRUE)
yhat.bag <- predict(rf.data, newdata=data.test)
plot(yhat.bag, meta.test)
table(yhat.bag, meta.test)
(19+24)/62
importance(rf.data)

> table(yhat.bag, meta.test)
      meta.test
yhat.bag DM  NODM
      DM   18   10
      NODM 10   24
> (18+24)/62
[1] 0.6774194
```

1.2.7 Boosting

```
## converting to binary (0,1) response
data.gbm <- data.lasso[, -1]
boolean <- ifelse(data.lasso$meta=="NODM", 1, 0)
data.gbm <- data.frame(boolean, data.gbm)
data.test <- data.gbm[-train,]

boost.data <- gbm(boolean~., data=data.gbm[train,], distribution="bernoulli",
                  n.trees=5000, interaction.depth=4)

yhat.boosting <- predict(boost.data, newdata=data.test, n.trees=100, distribution="bernoulli")
yhat.pred <- rep(0, 62)
```



```

yhat.pred[yhat.boosting>.5]=1
plot(yhat.boosting, data.test)
table(yhat.pred, boolean[-train])
(27+12)/62

> table(yhat.pred, boolean[-train])

yhat.pred  0  1
           0 27 22
           1  1 12
> (27+12)/62
[1] 0.6290323

```

1.2.8 Importance comparisons

As can be observed from the importance summary below, there isn't much difference between random forest and bagging in terms of importance of predictors. The order of the predictors in the case of boosting was quite different.

```

pdf("bag_data.pdf")
varImpPlot(bag.data)
dev.off()

```

```

pdf("rf_data.pdf")
varImpPlot(rf.data)
dev.off()

```

```

pdf("boost_data.pdf")
summary(boost.data)
dev.off()

```

```

> importance(bag.data)

```

	DM	NODM	MeanDecreaseAccuracy	MeanDecreaseGini
NM_000918	3.5037728	8.781774	9.182071	4.522524
NM_003258	7.1786980	6.169646	8.930712	6.413021
NM_004119	3.9275676	4.869651	6.053791	3.988813
AF279865	15.0288601	3.615650	13.997609	7.107205
NM_002811	12.8839394	5.005679	12.620362	6.377391
Contig48919_RC	7.8892269	8.526424	11.517478	6.630566
NM_003714	7.3136271	8.657350	10.483121	5.076347
NM_006054	3.3456481	5.128653	5.631485	3.451874
NM_007267	7.3468067	7.948700	10.334150	4.822468
AF055033	7.6168989	5.154784	8.312727	4.446796
Contig59134_RC	-0.5793959	5.179788	3.915095	2.121363
NM_001007	10.9758645	4.700569	10.242585	4.866076

```

> importance(rf.data)

```


	DM	NODM	MeanDecreaseAccuracy	MeanDecreaseGini
--	----	------	----------------------	------------------

NM_000918	2.958896	8.473358	7.940691	4.716304
NM_003258	7.586545	7.203475	10.225379	6.139778
NM_004119	4.718333	4.904367	6.909594	4.366215
AF279865	12.520495	4.814431	11.358147	6.170585
NM_002811	8.764961	6.550887	10.304826	6.180394
Contig48919_RC	6.635444	7.277151	9.369255	5.369459
NM_003714	5.290677	5.371936	6.926086	4.734545
NM_006054	4.535794	3.973607	5.933265	3.849397
NM_007267	6.576916	7.798570	9.665736	5.724488
AF055033	8.064672	2.530757	6.983284	4.409412
Contig59134_RC	2.871751	4.405412	4.903617	3.420377
NM_001007	7.976717	6.581179	9.594289	4.901141


```
> summary(boost.data)
```


	var	rel.inf
AF279865	AF279865	11.531043
NM_002811	NM_002811	10.949761
Contig48919_RC	Contig48919_RC	10.296364
NM_007267	NM_007267	9.600293
NM_006054	NM_006054	8.549060
NM_001007	NM_001007	8.182846
NM_003258	NM_003258	8.075747
AF055033	AF055033	7.537565
NM_003714	NM_003714	7.276408
NM_000918	NM_000918	6.789771
NM_004119	NM_004119	6.271262
Contig59134_RC	Contig59134_RC	4.939880

bag.data



rf.data





1.3 Annex

