

# Statistical methods for bioinformatics

## GAM and Trees

Cedric Lood

May 14, 2016

## 1 Applied exercises

### 1.1 Question 10

Libraries and definition of the training and test sets used for the analysis:

```
library(ISLR)
library(boot)
library(ggplot2)
library(leaps)
library(gam)

attach(College)
set.seed(1)
train <- sample(c(TRUE,FALSE), nrow(College), rep=TRUE)
test <- (!train)
```

#### 1.1.1 Part a

The analysis reveals a forward selection with 6 variables retained in the model:

```
model.fwd <- regsubsets(Outstate~.,data=College[train,], nvmax=17,method="forward")
test.mat <- model.matrix(Outstate~., data=College[test,])
```

```
val.errors <- rep(NA, 17)
for(i in 1:17){
  coefi <- coef(model.fwd, id=i)
  pred <- test.mat[,names(coefi)] %*% coefi
  val.errors[i] <- mean((College$Outstate[test]-pred)^2)
}
val.errors
which.min(val.errors)
coef(model.fwd, 6)
```

```
> val.errors
[1] 10734659 7647452 6468424 5434378 4948421 4650921 4734848 4718857
[9] 4690641 4735896 4694172 4877469 4754741 4690085 4711840 4700313
```

```
[17] 4700869
> which.min(val.errors)
[1] 6
> coef(model.fwd, 6)
      (Intercept)      PrivateYes      Room.Board      Terminal      perc.alumni
-4227.6797221    2778.7052614        0.8106532    49.5653264    38.3960635
      Expend      Grad.Rate
    0.2616141    26.3975500
```

### 1.1.2 Part b

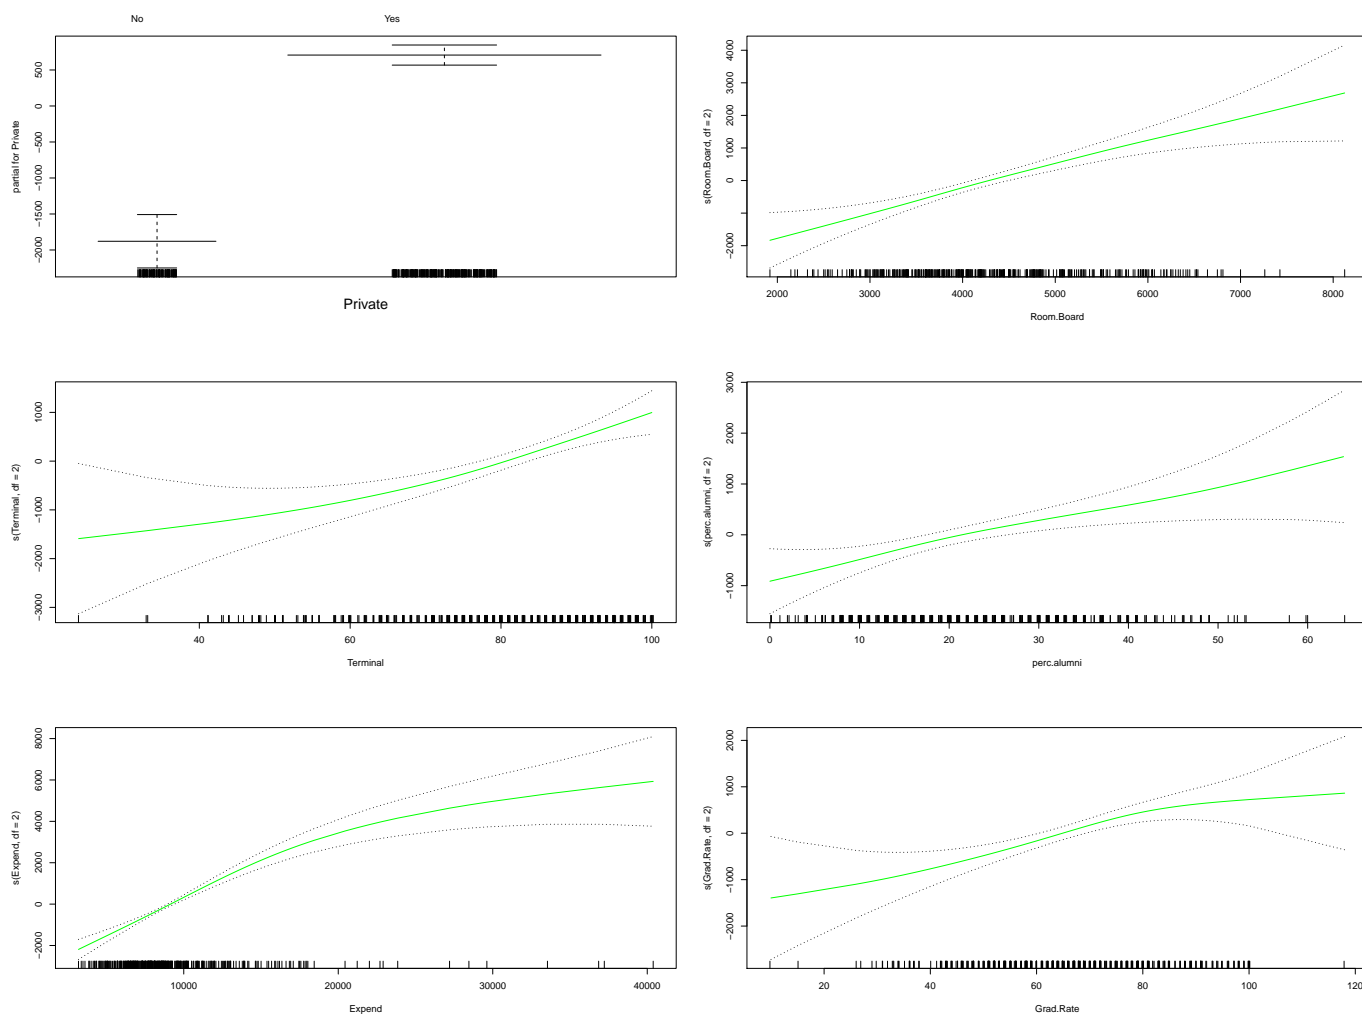
For this section, there are potentially many different GAM that could be produced. For example, one could produce different functionals consisting of smoothing splines for each of the predictors, with different combinations of degrees of freedom. Here are a few, the first one being akin to a multiple linear regression:

```
model.gam <- gam(Outstate~Private+Room.Board+Terminal+perc.alumni+Expend+Grad.Rate,data=Colleg
pdf("gam_trees_s.pdf", width=16, height=12)
par(mfrow = c(3, 2))
plot.gam(model.gam,se=TRUE,col="green")
dev.off()
```

```
model.gam.s2 <- gam(Outstate~Private+s(Room.Board,df=2)+s(Terminal,df=2)+s(perc.alumni,df=2)+s
pdf("gam_trees_s_2.pdf", width=16, height=12)
par(mfrow = c(3, 2))
plot.gam(model.gam.s,se=TRUE,col="green")
dev.off()
```

```
model.gam.s.3 <- gam(Outstate~Private+s(Room.Board,df=3)+s(Terminal,df=3)+s(perc.alumni,df=3)+
pdf("gam_trees_s_3.pdf", width=16, height=12)
par(mfrow = c(3, 2))
plot.gam(model.gam.s.3,se=TRUE,col="green")
dev.off()
```

The graphics for the other 2 models can be found in the annex. Here is the one for the smoothing splines with 2 degrees of freedom:



### 1.1.3 Part c

This is the evaluation of the 3 different models for the test sets, followed by the results.

```
gam.pred <- predict(model.gam,College[test,])
gam.err <- mean((College[test,]$Outstate - gam.pred)^2)

gam.pred.s2 <- predict(model.gam.s2,College[test,])
gam.err.s2 <- mean((College[test,]$Outstate - gam.pred.s2)^2)

gam.pred.s3 <- predict(model.gam.s.3,College[test,])
gam.err.s3 <- mean((College[test,]$Outstate - gam.pred.s3)^2)

> gam.err
[1] 4650921
> gam.err.s2
[1] 3952853
> gam.err.s3
[1] 3784922
```

#### 1.1.4 Part d

Evidence of non-linear relationship between variables can be investigated using the summary function on the GAM. The function provides the user with a non-parametric ANOVA table. Below is the output of the command, which seem to indicate a non-linear relationship between the predictor “Expend” and the dependent value.

```
> summary(model.gam.s2)
```

```
Call: gam(formula = Outstate ~ Private + s(Room.Board, df = 2) + s(Terminal,
      df = 2) + s(perc.alumni, df = 2) + s(Expend, df = 2) + s(Grad.Rate,
      df = 2), data = College[train, ])
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-7045.4	-1145.4	84.8	1184.1	5047.0

(Dispersion Parameter for gaussian family taken to be 3452101)

Null Deviance: 6006262152 on 405 degrees of freedom

Residual Deviance: 1360128366 on 394.0002 degrees of freedom

AIC: 7278.121

Number of Local Scoring Iterations: 2

Anova for Parametric Effects

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
Private	1	1635387248	1635387248	473.737	< 2.2e-16 ***
s(Room.Board, df = 2)	1	1343549886	1343549886	389.198	< 2.2e-16 ***
s(Terminal, df = 2)	1	597441375	597441375	173.066	< 2.2e-16 ***
s(perc.alumni, df = 2)	1	240771844	240771844	69.746	1.161e-15 ***
s(Expend, df = 2)	1	424246993	424246993	122.895	< 2.2e-16 ***
s(Grad.Rate, df = 2)	1	63996091	63996091	18.538	2.104e-05 ***
Residuals	394	1360128366	3452101		

---

Signif. codes: 0 ‘\*\*\*’ 0.001 ‘\*\*’ 0.01 ‘\*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Anova for Nonparametric Effects

	Npar	Df	Npar F	Pr(F)
(Intercept)				
Private				
s(Room.Board, df = 2)	1	0.5670	0.4519	
s(Terminal, df = 2)	1	2.2148	0.1375	
s(perc.alumni, df = 2)	1	1.0512	0.3059	
s(Expend, df = 2)	1	26.1955	4.837e-07	***
s(Grad.Rate, df = 2)	1	3.5803	0.0592	.

---

Signif. codes: 0 ‘\*\*\*’ 0.001 ‘\*\*’ 0.01 ‘\*’ 0.05 ‘.’ 0.1 ‘ ’ 1

## 1.2 Trees Vijver

### 1.2.1 Performance with Ridge and Lasso

For reminder, here are the performance obtained using regularization techniques:

```
> perf.ridge  
[1] 0.6489362
```

```
> perf.lasso  
[1] 0.6170213
```

### 1.3 Annex

